

KUNAL NISHAD

+91 9621088062 | [linkedin.com/in/kunal-nishad-643838250](https://www.linkedin.com/in/kunal-nishad-643838250) | nishadkunal1234567@gmail.com | github.com/kunalnishad115

EDUCATION

GLA University, Mathura

Aug 2022 – Aug 2026

B.Tech in Computer Science; CGPA: 7.6

EXPERIENCE

Data Engineer Intern — Ai Tech Tures Lab LLP

Aug 2026 – Present

Mumbai, India

- Assist the data engineering team in building and maintaining **ETL/ELT pipelines across AWS and Azure** to ingest and process the Midas Jewel dataset, supporting relevant analytics for the Data Analytics (DA) team.
- Support **data modeling and data design**, and lead the merging of Midas Jewel source data into the application database to ensure accurate, consistent data across systems.

PROJECTS

Production-Ready RAG Application with Hybrid Retrieval & Source Citations

Python, Streamlit, LangChain, ChromaDB, BM25, Cross-Encoder, Sentence Transformers, Gemini, Cohere, Docker

- Architected an end-to-end Retrieval-Augmented Generation (RAG) pipeline** covering document ingestion (PDF/DOCX/TXT), text cleaning, metadata extraction, chunking, and **embedding generation** with Sentence Transformers.
- Implemented hybrid retrieval combining ChromaDB vector search** with a BM25 keyword index, followed by cross-encoder re-ranking, improving retrieval accuracy and reducing hallucinated responses.
- Built an interactive Streamlit application** integrating Gemini and Cohere for LLM response generation, applying **prompt engineering** techniques to ground every answer in retrieved chunks and return inline source citations (document name and page/section) for traceability.

End-to-End Azure Data Engineering Pipeline for Amazon Review Analytics

Azure Data Factory, Databricks, ADLS Gen2, Azure Synapse Analytics, Power BI

- Built an end-to-end Azure pipeline (ADF, ADLS Gen2, Databricks, Synapse) using Medallion Architecture; developed scalable **PySpark ETL** to clean and validate **20K+ Amazon reviews across 123 products**, achieving **zero missing or duplicate records**.
- Designed business-ready Gold datasets and Synapse Serverless SQL views to generate KPIs and trend analysis, visualized through interactive Power BI dashboards.

Zomato AI Data Engineering Project

AWS S3, Snowflake, dbt, Airflow, LangChain, Groq, Cohere, Streamlit

- Built an **ELT pipeline on AWS S3 + Snowflake** using **dbt** (Medallion architecture) to transform a 10M+ row dataset, orchestrated via **Apache Airflow (Docker)**, with automated data-quality tests at every layer.
- Added an **LLM layer with LangChain (Groq + Cohere)** — enrichment for review sentiment, a **RAG app** for review search, and a **text-to-SQL** app for warehouse queries, all served via **Streamlit**.

CERTIFICATES / ACHIEVEMENTS

- Coding Weekly Rank 30 — Unstop.

TECHNICAL SKILLS

Programming Languages: Python, SQL

Big Data & Processing: PySpark, Pandas, NumPy

Cloud & Data Platforms: ADF, Azure Databricks, ADLS Gen2, Azure Synapse, AWS-S3, Snowflake, Delta Lake

Orchestration & DevOps: Apache Airflow, dbt, Docker, Git

Databases: MySQL, PostgreSQL, ChromaDB, Redis

GenAI/LLM: Lang Chain, Lang Graph, Chroma DB, Streamlit